



# Optimization with Moose

# General optimization tips

# The first rule of optimization is:

# The first rule of optimization is: don't!

Is your time  
better spent on  
features?  
Bug fixes?

*"There are only two hard problems in Computer Science: cache invalidation and naming things."*

# Optimization often comes at the expense of maintainability

# You're expensive!



# The first rule of optimization is: don't!

# The second rule of optimization: profile!

# Don't assume

# Be scientific

- Start by profiling
- Make changes
- Profile again!

# Modern computers are complex

- ▶ Instruction pipelining
- ▶ CPU cache lines
- ▶ Filesystem cache
- ▶ OS process and thread scheduling

# Profile *real* workloads

# One-time and amortized costs

# Be wary of micro- benchmarks



# You workload changes

# Your code changes

# The second rule of optimization: profile!

# Profilers

- Devel::NYTProf
- DTrace

# Devel::NYTProf

- ▶ Perl-specific
- ▶ Line-, sub-, block-, and opcode-level profiling
- ▶ Low overhead
- ▶ Accurate

# Live demo!

- `cd Moose`
- `export PERL50PT=-d:NYTProf`
- `ack Moose >/dev/null`
- `export PERL50PT=`
- `nytprofhtml --open`

# No demo?

## Performance Profile Index

For /Users/sartak/.perl/perls/perl-5.16.0/bin/ack

Profile of /Users/sartak/.perl/perls/perl-5.16.0/bin/ack for 1.45s (of 1.82s), executing 846546 stat source files and 1 string evals.

### Top 15 Subroutines

Calls	P	F	Exclusive Time	Inclusive Time	Subroutine
536	1	1	473ms	1.23s	App::Ack:: <a href="#">search_resource</a>
78937	2	1	289ms	331ms	App::Ack::Resource::Basic:: <a href="#">next_text</a>
5157	1	1	94.1ms	94.1ms	App::Ack::Resource::Basic:: <a href="#">CORE:ftbinary</a> (opcode)
5157	1	1	90.0ms	215ms	App::Ack:: <a href="#">print_match_or_context</a>
79620	5	1	67.1ms	67.1ms	App::Ack:: <a href="#">CORE:regcomp</a> (opcode)
5838	2	1	49.6ms	49.6ms	App::Ack:: <a href="#">CORE:subst</a> (opcode)
15471	3	1	44.4ms	51.0ms	App::Ack:: <a href="#">print</a>
78937	1	1	42.0ms	42.0ms	App::Ack::Resource::Basic:: <a href="#">CORE:readline</a> (opcode)
83872	13	1	27.7ms	27.7ms	App::Ack:: <a href="#">CORE:match</a> (opcode)
5157	1	1	25.2ms	119ms	App::Ack::Resource::Basic:: <a href="#">is_binary</a>
1	1	1	20.6ms	1.39s	App::Ack:: <a href="#">print_matches</a>
673	1	1	20.6ms	20.6ms	App::Ack::Resource::Basic:: <a href="#">CORE:sysread</a> (opcode)
5157	1	1	12.9ms	27.9ms	App::Ack:: <a href="#">print_line_no</a>
1	1	1	12.8ms	12.8ms	utf8:: <a href="#">SWASHNEW</a>
5157	1	1	12.8ms	31.1ms	App::Ack:: <a href="#">print_filename</a>

See [all 383 subroutines](#)



# No demo?

2106	536	60µs			my \$has_lines = 0;
2107	536	57µs			my @lines;
2108	536	236µs			if ( defined \$opt->{lines} ) {
2109					\$has_lines = 1;
2110					@lines = ( @{\$opt->{lines}}, -1 );
2111					undef \$regex; # Don't match when printing matching line
2112					}
2113					else {
2114	536	2.95ms	1072	1.14ms	\$regex = qr/\$opt->{regex}/;
					# spent 986µs making 536 calls to App::Ack::CORE:qr, avg 2µs/call
					# spent 158µs making 536 calls to App::Ack::CORE:regcomp, avg 295ns/
2115					}
2116					
2117					# for context processing
2118	536	62µs			\$last_output_line = -1;
2119	536	43µs			\$any_output = 0;
2120	536	160µs			my \$before_context = \$opt->{before_context};
2121	536	67µs			my \$after_context = \$opt->{after_context};
2122					
2123	536	86µs			\$keep_context = (\$before_context    \$after_context) && !\$passthru;
2124					
2125	536	30µs			my @before;
2126	536	11µs			my \$before_starts_at_line;
2127	536	56µs			my \$after = 0; # number of lines still to print after a match
2128					
2129	536	78.0ms	73780	306ms	while ( \$res->next_text ) {
					# spent 306ms making 73780 calls to App::Ack::Resource::Basic::next_text
2130					# XXX Optimize away the case when there are no more @lines to find.
2131					# XXX \$has_lines, \$passthru and \$v never change. Optimize.
2132	78401	298ms	156802	91.7ms	if ( \$has_lines
					# spent 66.5ms making 78401 calls to App::Ack::CORE:regcomp, avg 849r
					# spent 25.2ms making 78401 calls to App::Ack::CORE:match, avg 321ns/
2133					? \$. != \$lines[0] # \$lines[0] should be a scalar
2134					: Sv ? m/\$regex/ : !m/\$regex/ \ }

# DTrace

- ▶ Not Perl-specific
- ▶ System profiler
- ▶ Profile your kernel too!
- ▶ Solaris, OS X, FreeBSD
- ▶ Linux support iffy
- ▶ Low overhead
- ▶ Production safe!



# DTrace

- ▶ Capture many kinds of events
- ▶ syscall
- ▶ memory allocation
- ▶ thread scheduling
- ▶ process lifecycle
- ▶ Perl function call
- ▶ Perl global phase change (BEGIN, END, etc)

# Live demo!

- ▶ How much time does each of ack's function calls spend in syscalls?
- ▶ 

```
sudo dtrace -qZn 'perl::sub-entry /substr(copyinstr(arg3), 0, 8) == "App::Ack"/ { self->wanted = pid; self->profiling = strjoin(copyinstr(arg3), strjoin("::", copyinstr(arg0))) } syscall::entry /self->wanted == pid/ { self->started = timestamp } syscall::return /self->started/ { @syscalls[self->profiling] = sum((timestamp - self->started) / 1000); self->started = 0 } proc::exit /self->wanted == pid/ { exit(0) }'
```
- ▶ `ack Moose >/dev/null`

# No demo?

App::Ack::print	11
App::Ack::read_ackrc	38
App::Ack::exit_from_ack	61
App::Ack::print_matches	63
App::Ack::Resource::Basic::reset	250
App::Ack::Resource::Basic::next_text	470
App::Ack::Resource::Basic::close	774
App::Ack::BEGIN	1124
App::Ack::Resource::Basic::needs_line_scan	1893
App::Ack::Repository::Basic::close	3077
App::Ack::Resource::Basic::new	3248
App::Ack::ignoredir_filter	3809
App::Ack::is_searchable	9581
App::Ack::Resource::Basic::is_binary	10877

# General optimizations

# Algorithm and data structure changes

# Array vs Hash

*"Doing linear scans over an associative array is like trying to club someone to death with a loaded Uzi."*

- Larry Wall

# Big-oh complexity

- ▶  $O(n)$
- ▶  $O(n^2)$
- ▶  $O(n^3)$
- ▶  $O(\log n)$
- ▶  $O(n \log n)$

# CPAN modules

- Set::Object
- Heap::Fibonacci
- Graph::Implicit
- Algorithm::



# Cache?

- use Memoized; memoize 'fib';
- Memcached

# Moose optimizations

# Moose is pretty fast

# Moose gives you knobs

# make\_immutable

```
__PACKAGE__->meta->make_immutable;
```

# make\_immutable

- ▶ string evals a constructor (“new”)
- ▶ and a destructor (“DESTRUCTALL”)
- ▶ memoizes meta-object methods

# string evals a constructor ("new")



string evals a  
destructor  
DEMOLISHALL

# Memoizes meta-object methods

# Use `make_immutable`

Use  
`make_immutable`  
except when  
you shouldn't!

*"There ain't no  
such thing as a  
free lunch"*

# Amortized cost

# Profile real workloads!

# Be kind: Leave a note



# Use `make_immutable` by default

# Sorry...

# Attributes vs Methods

# Memory usage

# Lazy

# “get” isn’t free

- is there a @\_?
- pull the value out of the object
  - (which is a hash lookup)
- laziness slows this down further

# Attributes aren't bad...

# Maybe use a method?



Does it really  
need to be  
stored for each  
object?

# Class constant

```
has some_value => (  
  is          => 'ro',  
  default => sub { 100 },  
);
```

```
sub some_value { 100 }
```

# Method instead of default/ builder?

# Role application

# Role application

≈

# inlining

# Trust me, I benchmarked it

Class::WithInheritance:  
1882648/s

Class::WithRoles:  
1923234/s

Class::WithInheritance  
567344/s

Class::WithRoles  
659223/s



Class::WithInheritance::AndClassNamesReallySeemToMatter  
434438/s

Class::WithRoles  
643734/s

# Lazy

```
has dbh => (  
    is      => 'ro',  
    default => sub { DBI->connect(...) },  
    lazy    => 1,  
);
```

# Lazy saves time and memory

*"There ain't no  
such thing as a  
free lunch"*

# Lazy's overhead: the accessor

# “get” isn’t free

- ▶ is there a @\_?
- ▶ pull the value out of the object
  - ▶ (which is a hash lookup)
- ▶ laziness slows this down further

# Accessors - lazy

- ▶ Moose blindly returns the attribute's value

# Accessors + lazy

- ▶ object, do you have a value for this attribute?
  - ▶ `exists $self->{$attribute}`
- ▶ if so, pull it out and return it
- ▶ if not:
  - ▶ invoke the default/builder
  - ▶ set the value
    - ▶ which checks type, might run coercions
- ▶ return it



Don't use lazy  
blindly  
everywhere

# Use a method?

Measure!  
Profile!  
Don't assume!

# Detour: lazy's other function

# Consulting other attributes during initialization

# Lazy's functions

- ▶ deferring expensive attribute initialization
- ▶ consulting other attributes during initialization

# Type constraints

# Type coercions

# Defining types

```
use Moose::Util::TypeConstraints;  
subtype 'Price',  
    as 'Str',  
    where {  
        /^ \p{Currency_Symbol} \d+ (\.\d+)? $/x  
    };
```



# Defining types

```
package Car;  
use Moose;  
has sticker_price => (  
    is => 'rw',  
    isa => 'Price',  
);
```

# Tricky types

- ▶ check as many aspects of the value as you want
- ▶ database lookups
- ▶ network requests?
- ▶ Mechanical Turk?
- ▶ arbitrary code

# subtype ‘Price’

- ▶ number of decimal digits (\$ £ €, ¥ , gas prices)
- ▶ optional commas in the right places
- ▶ localization (e.g. “1.000.000,00”)

# Expensive!

# Data normalization?

```
package Car;
use Moose;

has sticker_price => (
    is => 'rw',
    isa => 'Int',
);

sub formatted_price { ... }
```

# Boundary types

# Type coercion can be expensive too



# Type coercion

- ▶ Involves many type constraint checks
  - ▶ Check each potential “from” constraint
  - ▶ Run the coercion’s transform code
  - ▶ Ensure the output matches the “to” constraint

# Type coercion

- ▶ Makes your API more complex
  - ▶ “I can accept a string or an arrayref or a regular expression or a DateTime!”
- ▶ You can afford to be strict in your APIs

# Avoid coercions

- or, use them only at the boundary of your API

# Types are expensive

# But don't discard them

# “Not a HASH reference”

# Type inheritance and `inline_as`

# Type validation

- ▶ Types have hierarchy
- ▶ Each type (except the top-level “Any”) has a supertype
- ▶ Each type *specializes* its supertype




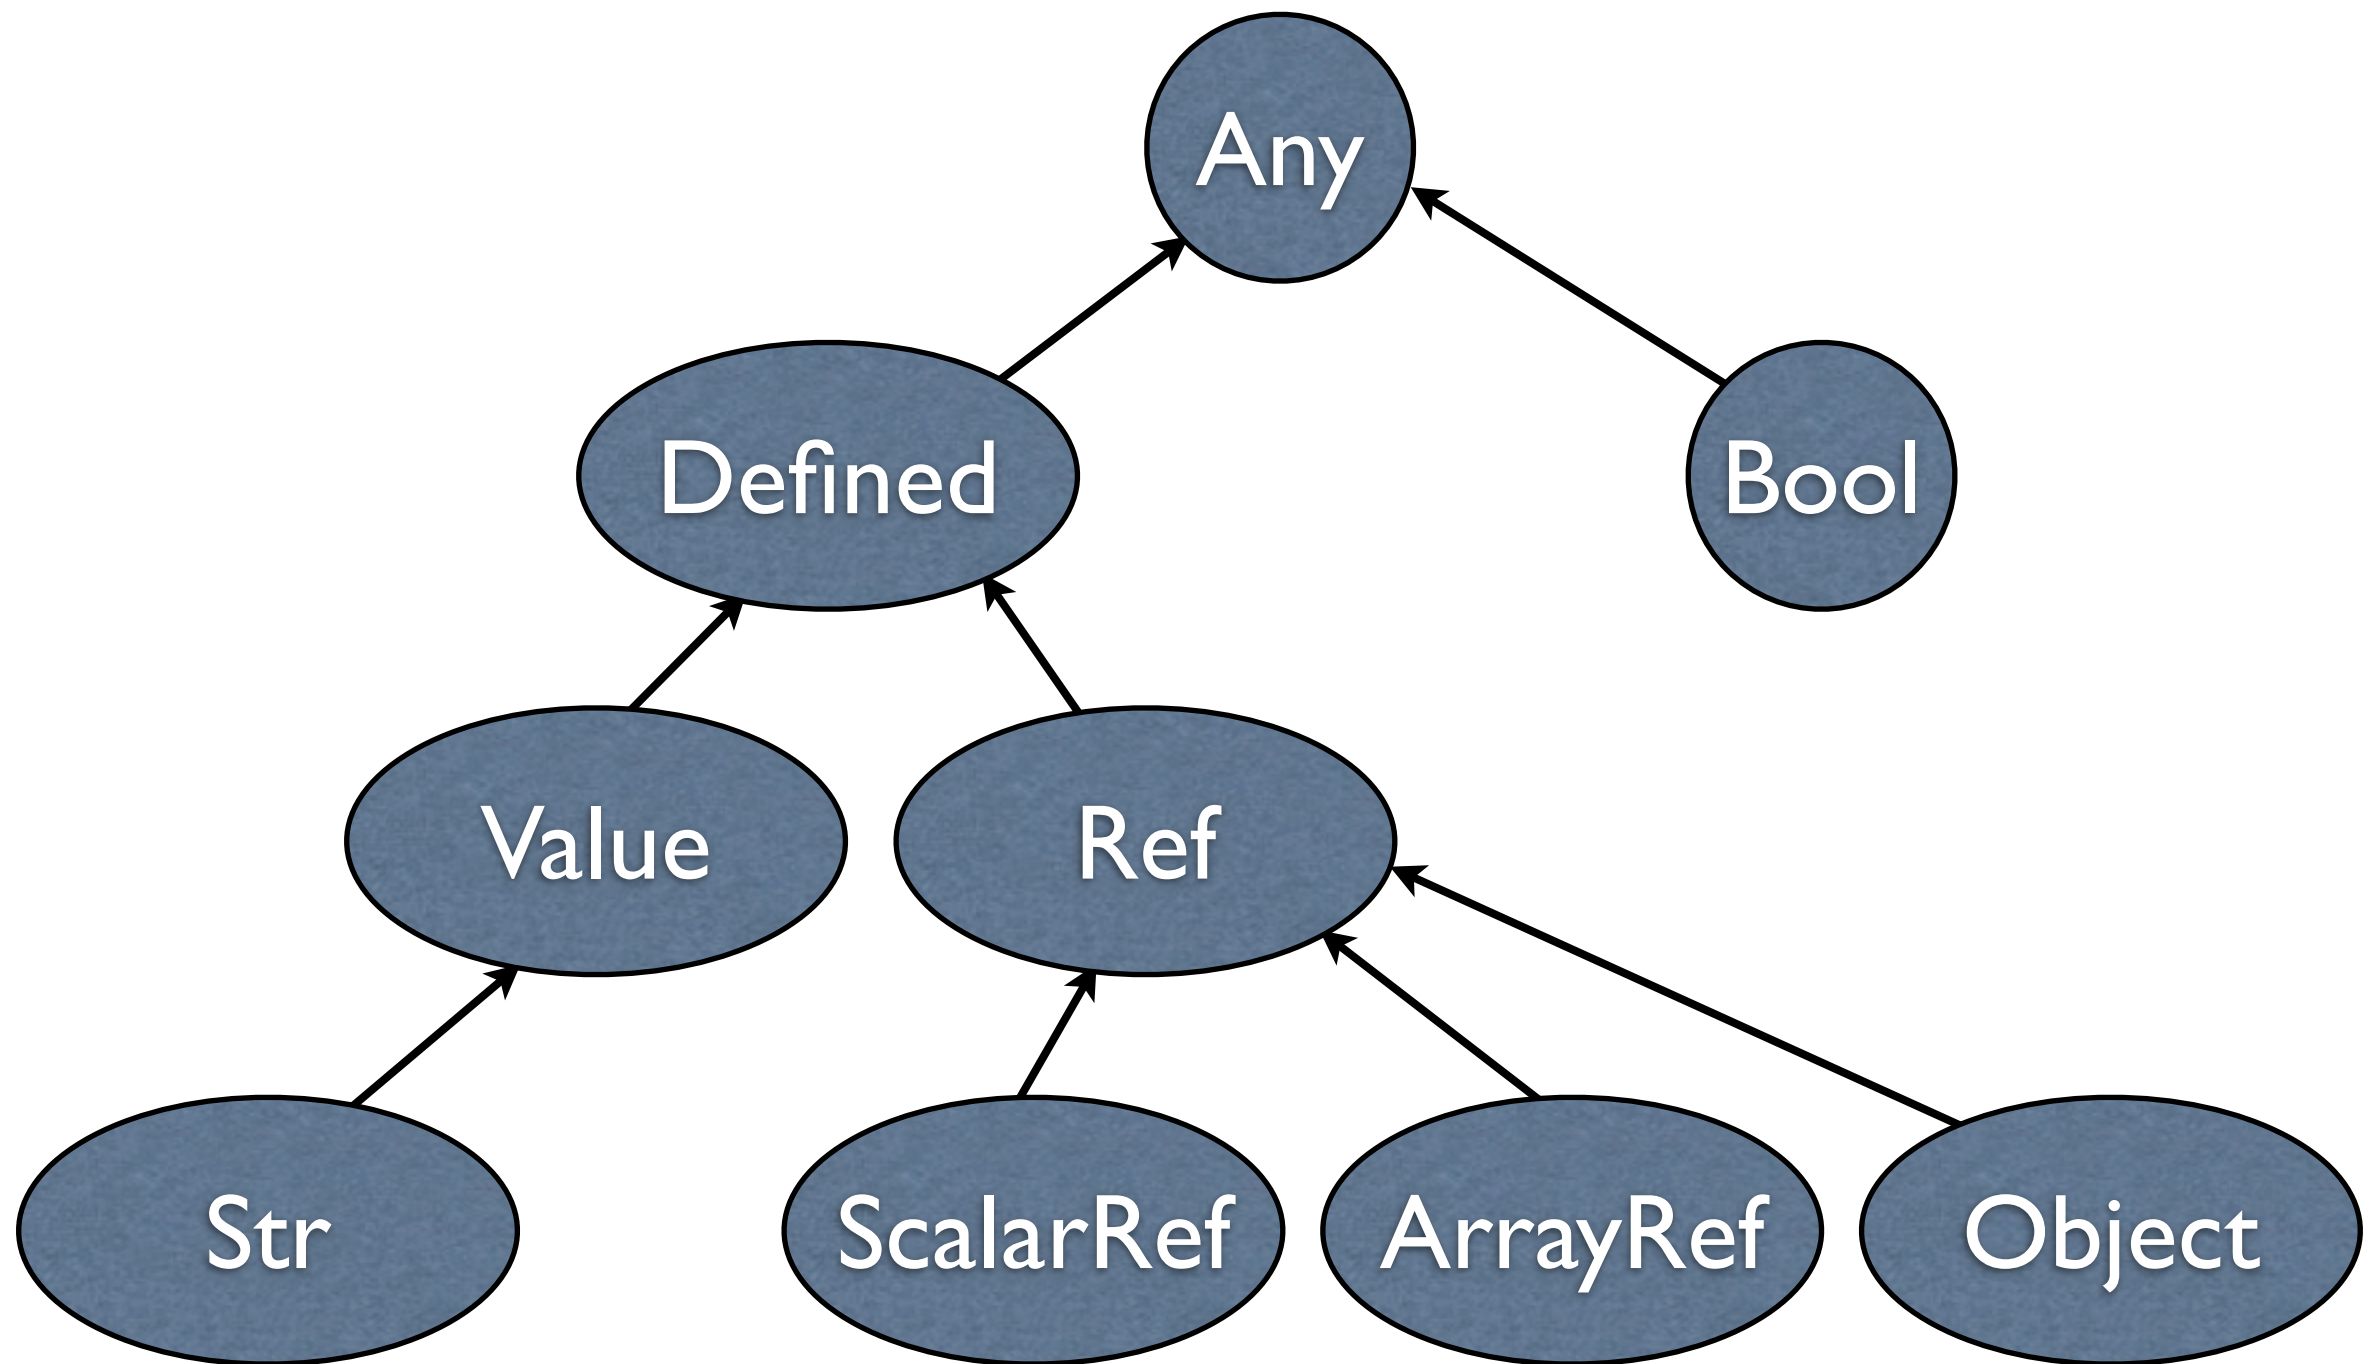
# Type validation

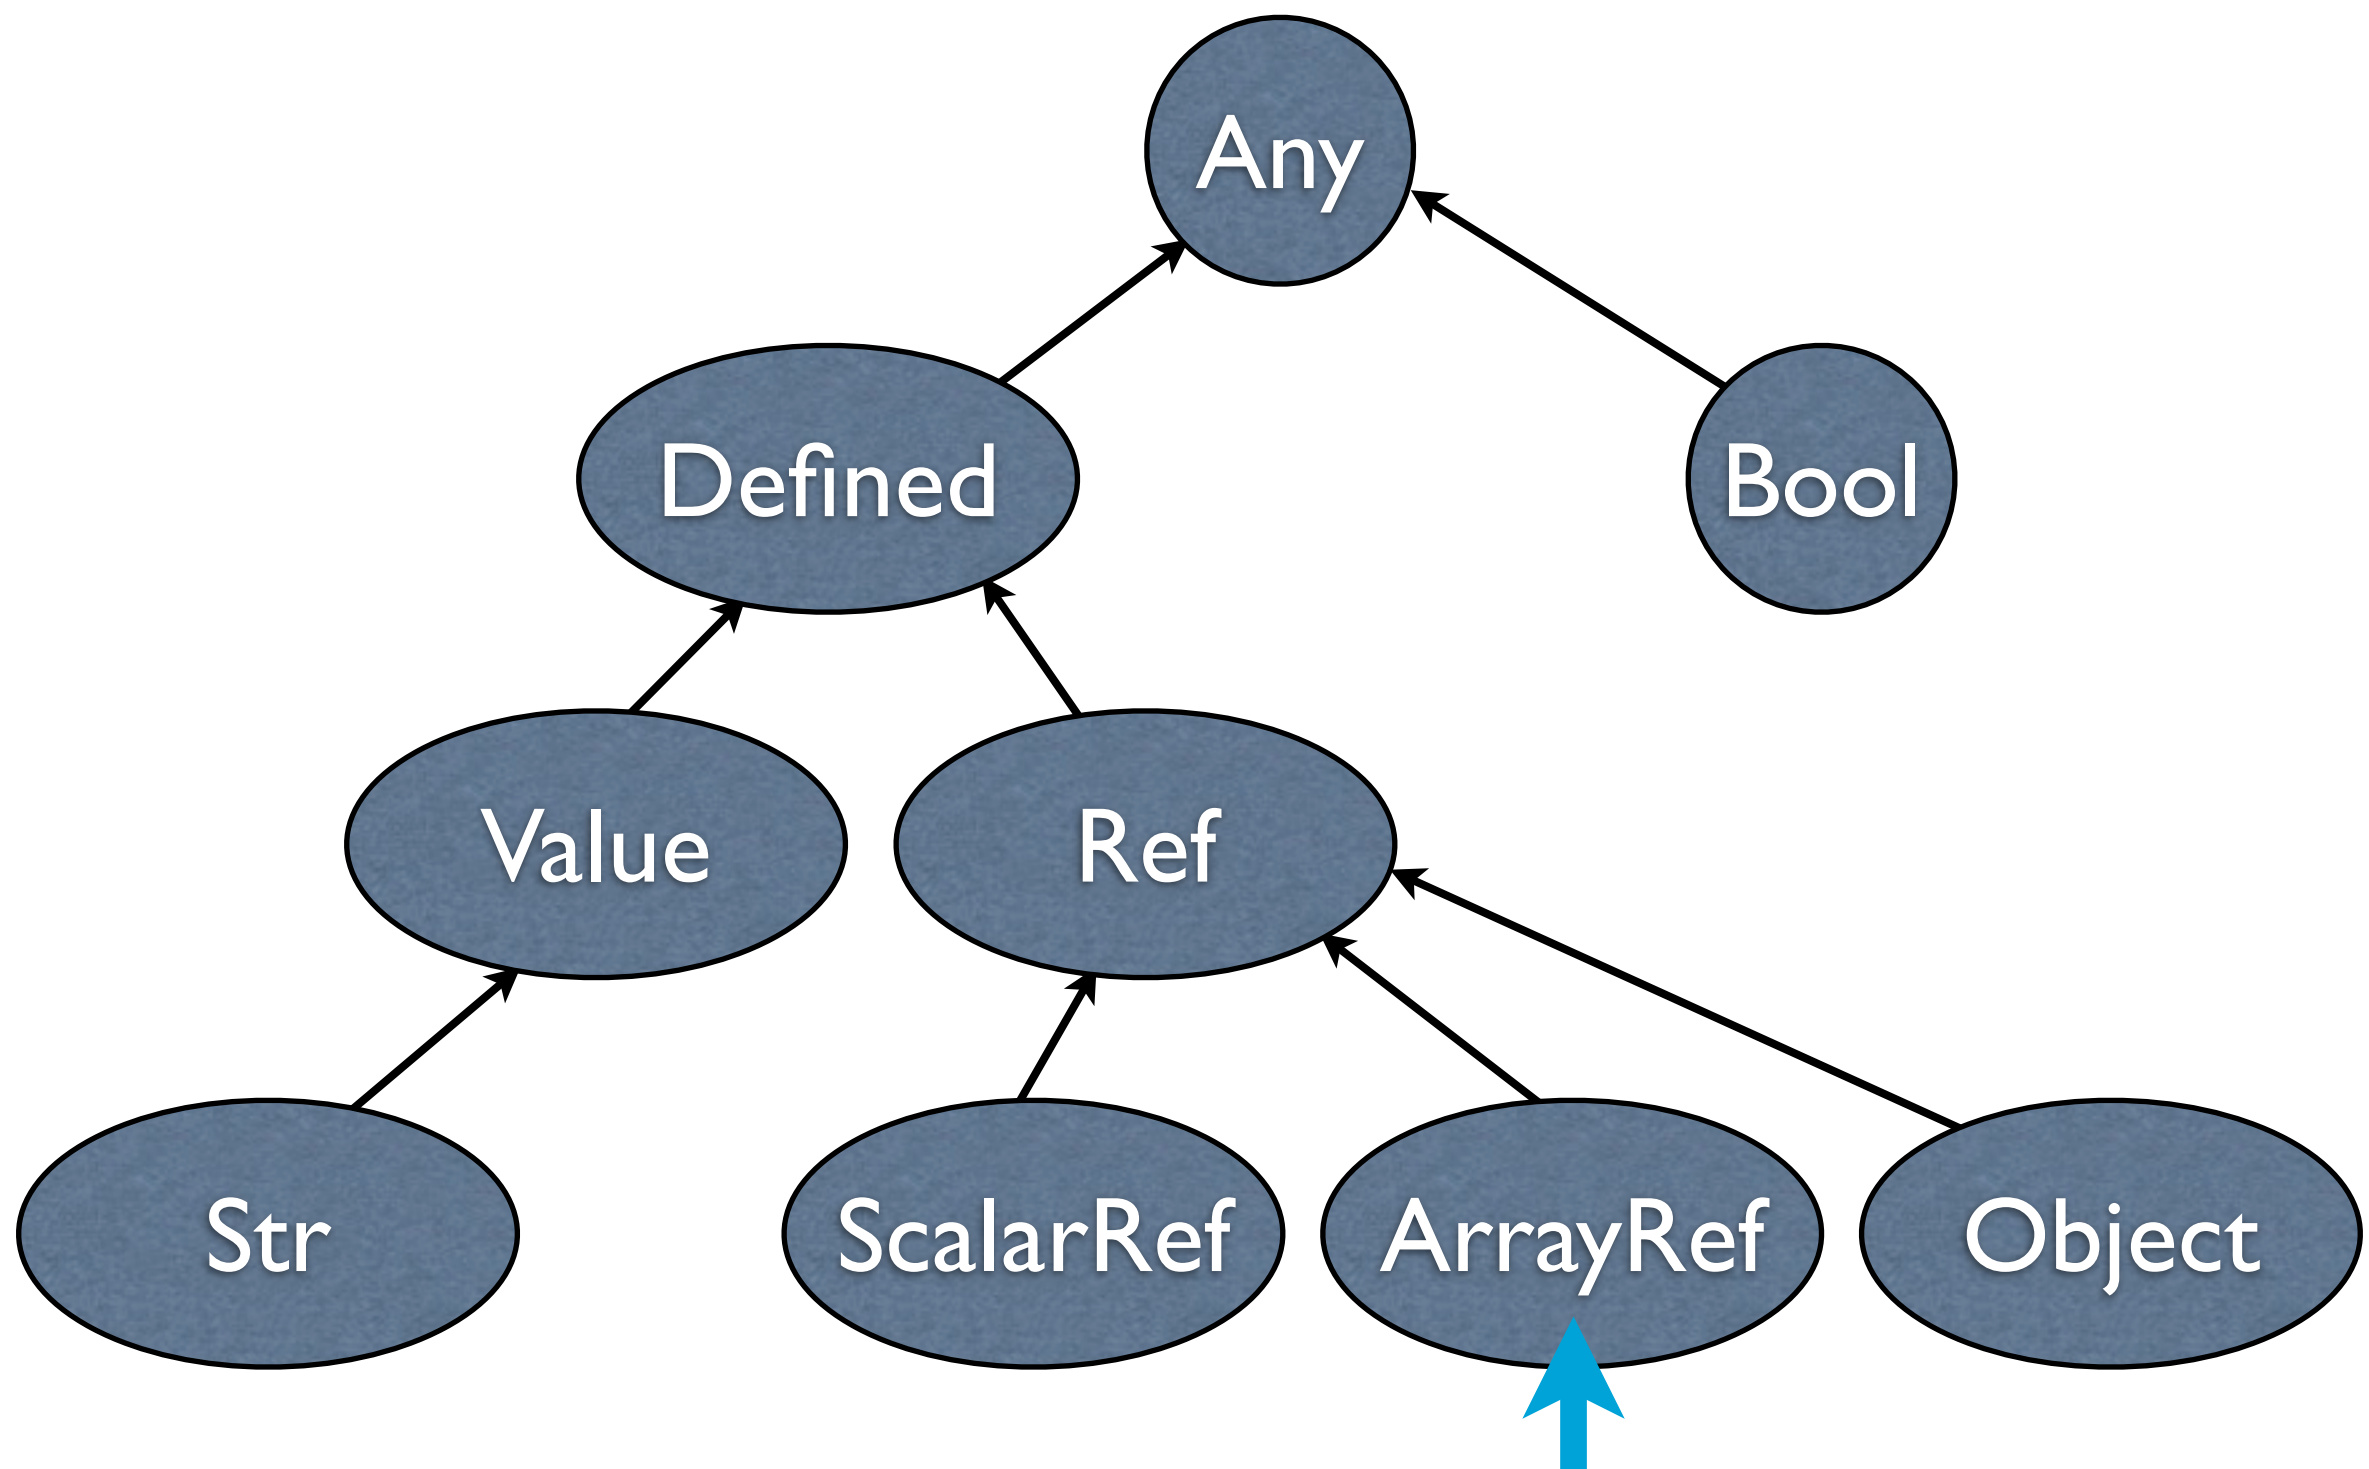
- Types have hierarchy
- Each type (except the top-level “Any”) has a supertype
- Each type *specializes* its supertype
- Each type check first validates against the supertype’s constraint, recursively

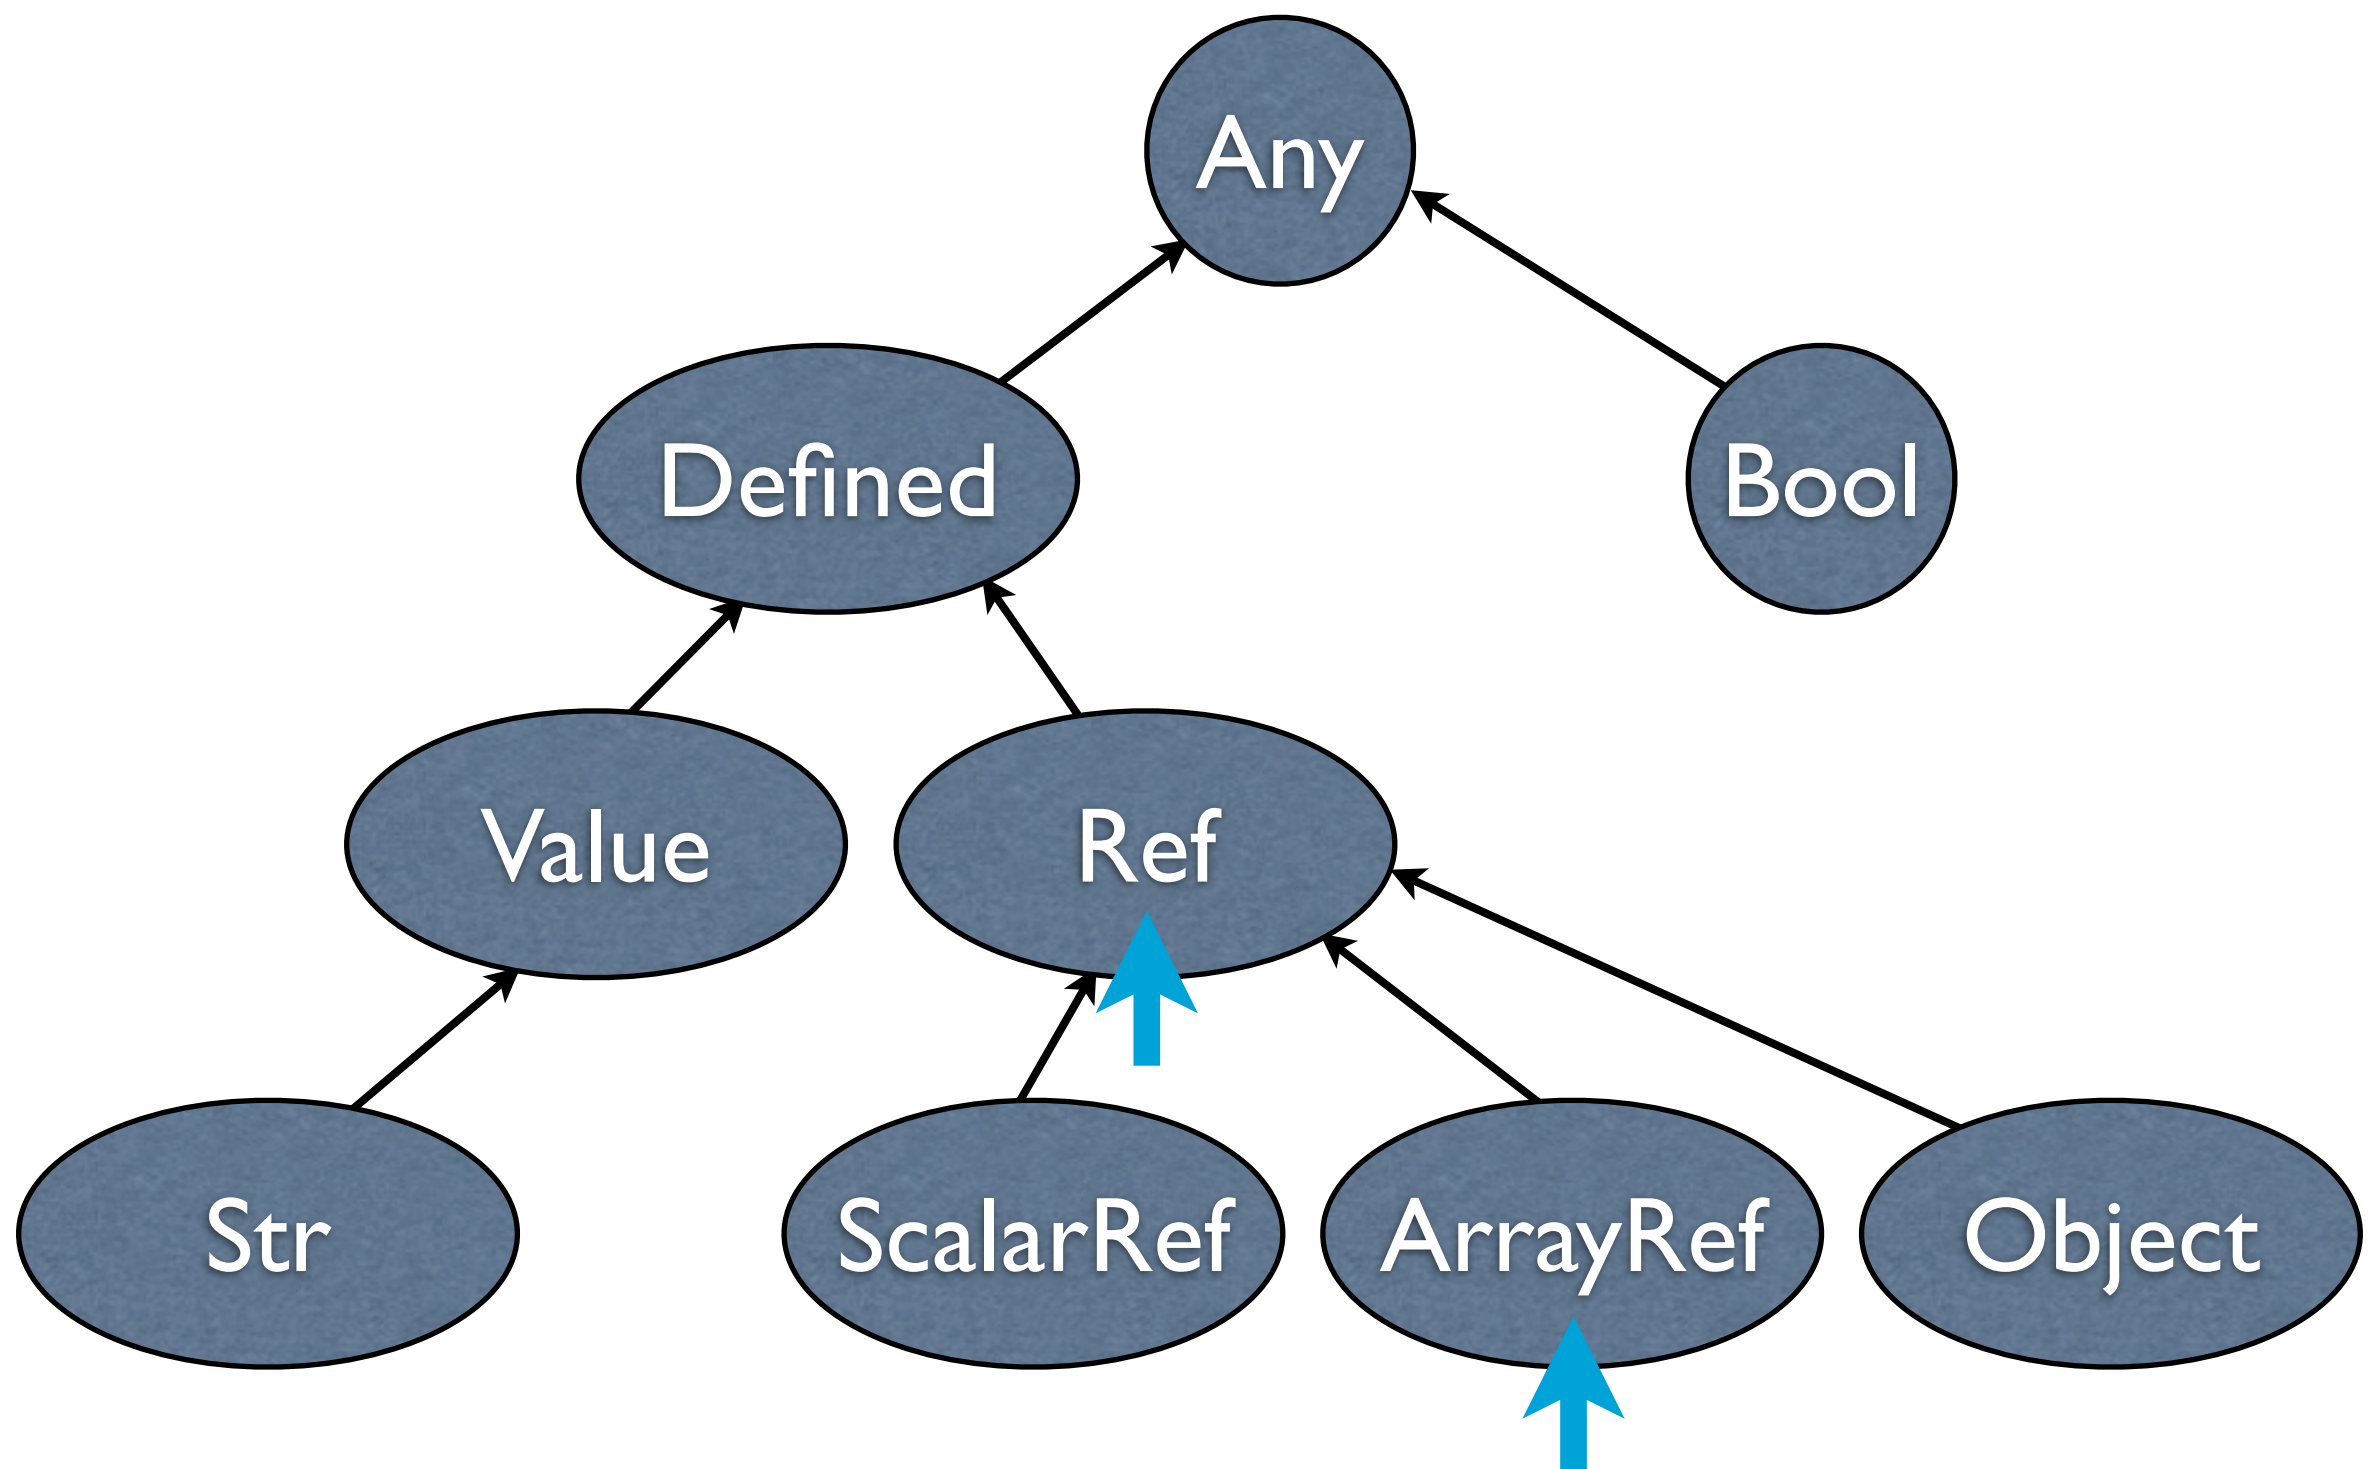
# Type validation

```
subtype 'PositiveInt'  
  as 'Int',  
  where { $_ > 0 };
```

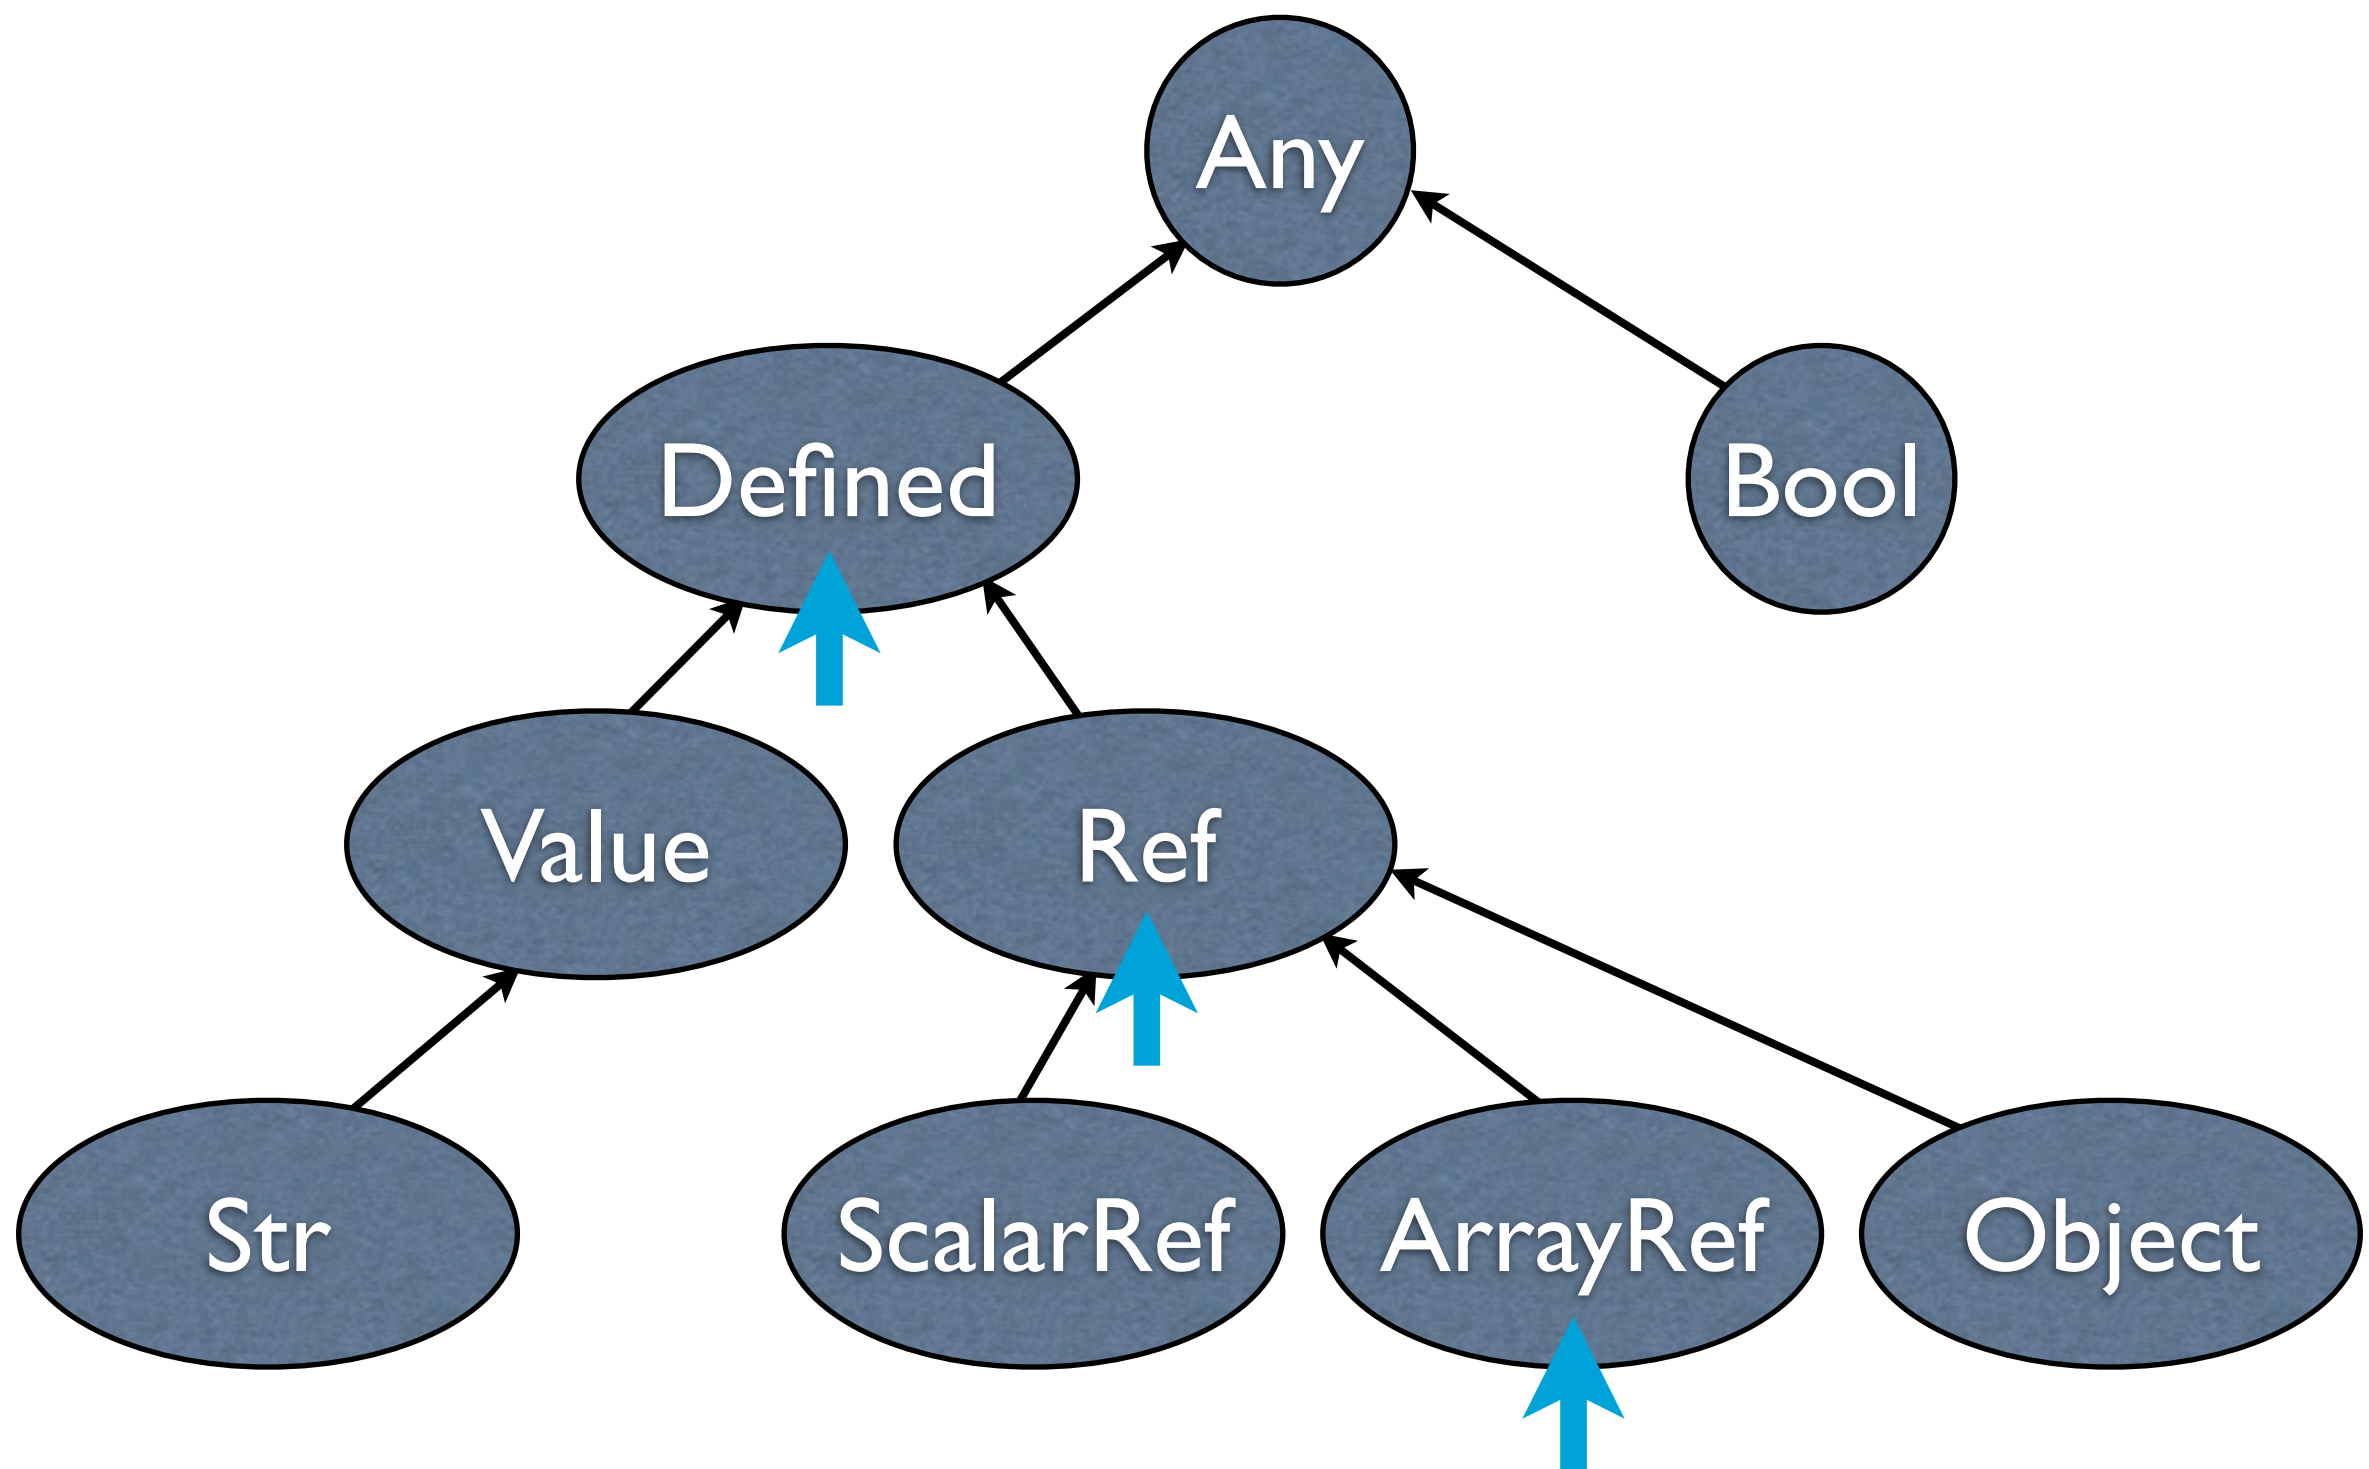


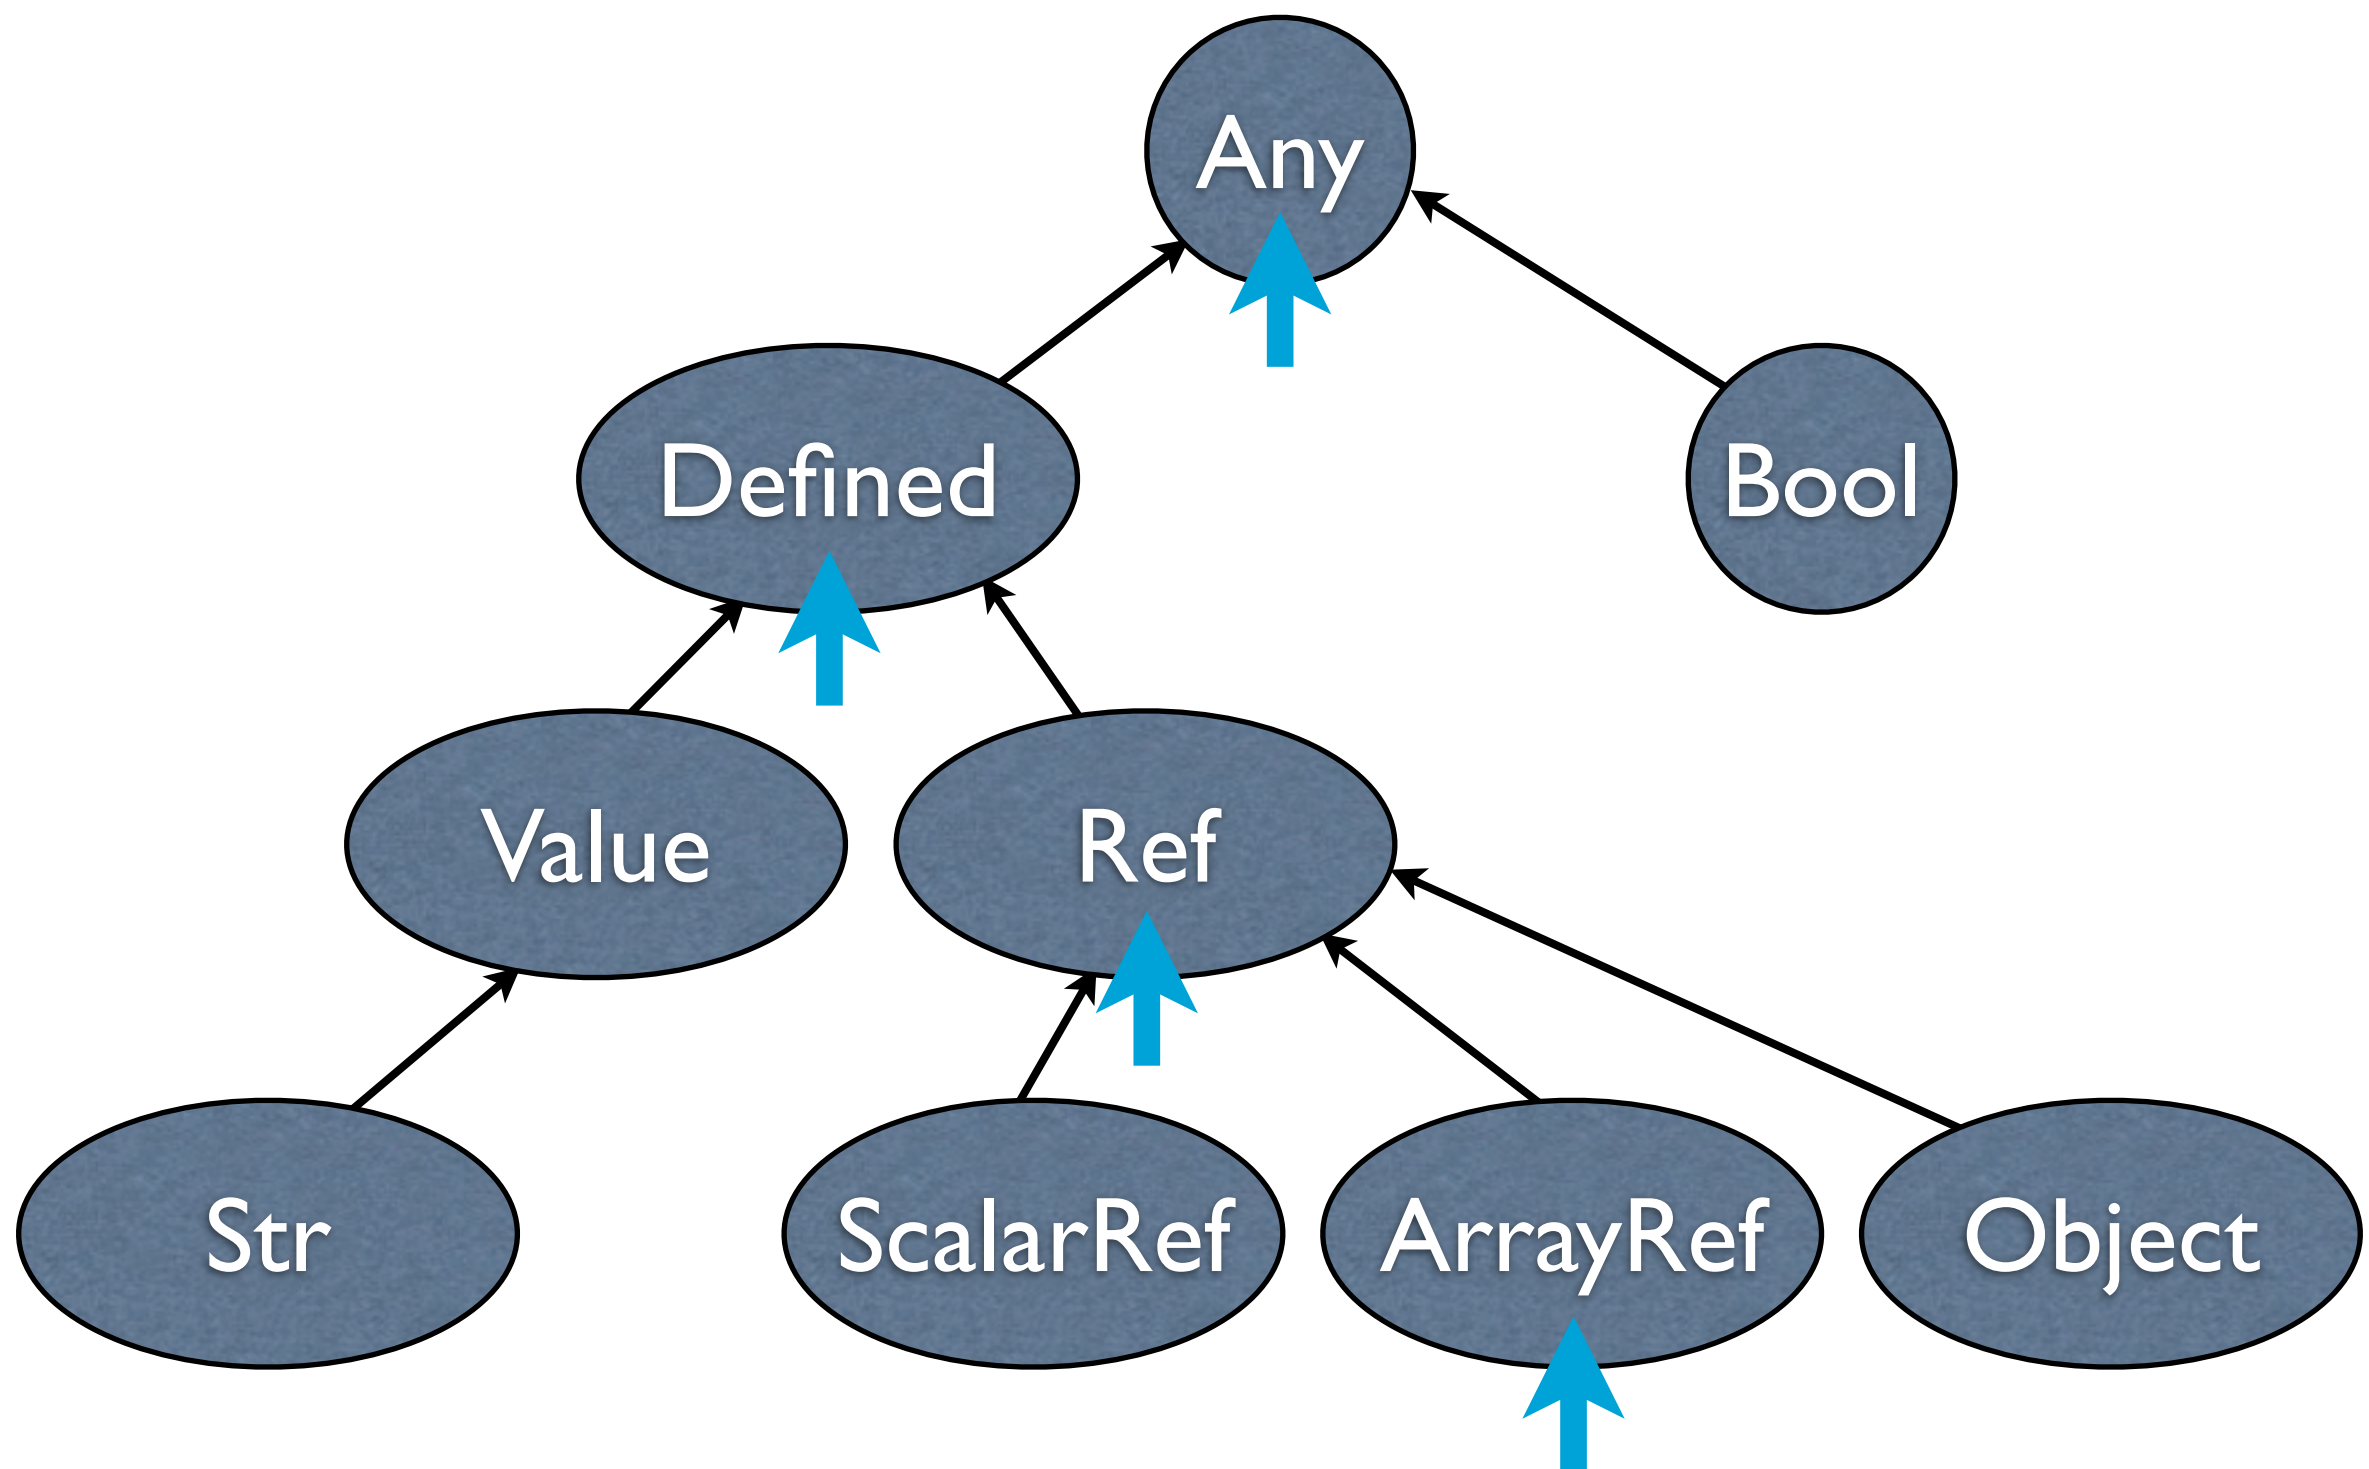




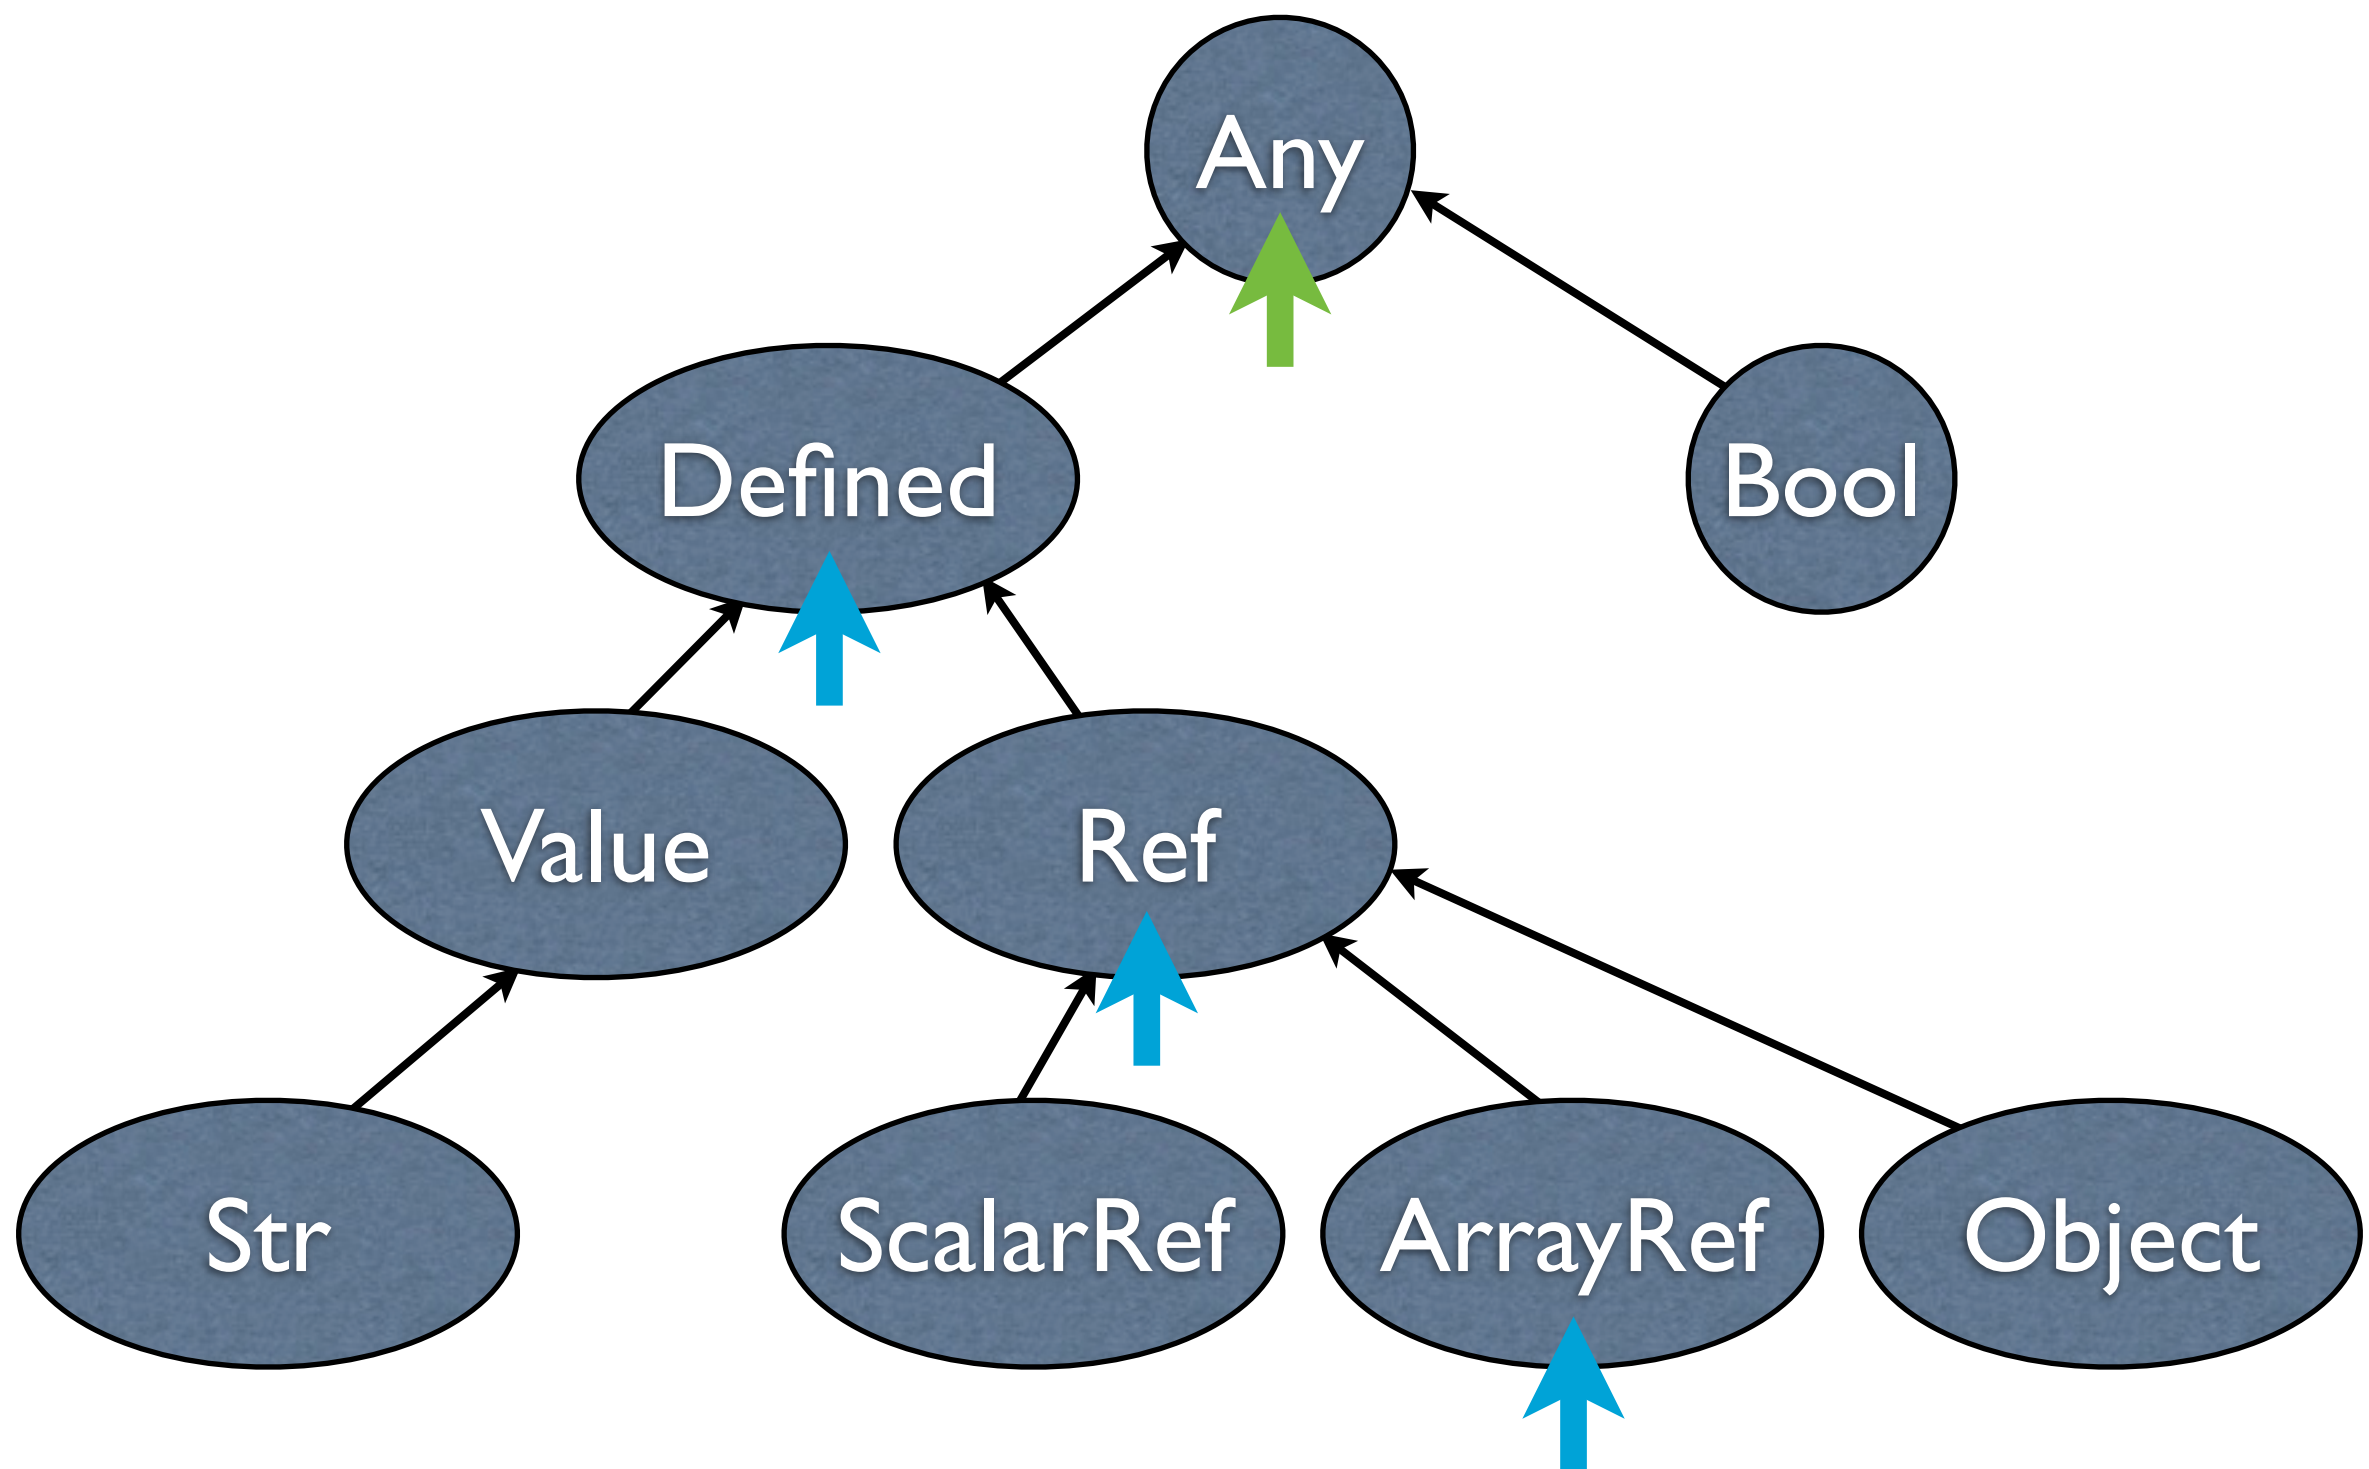


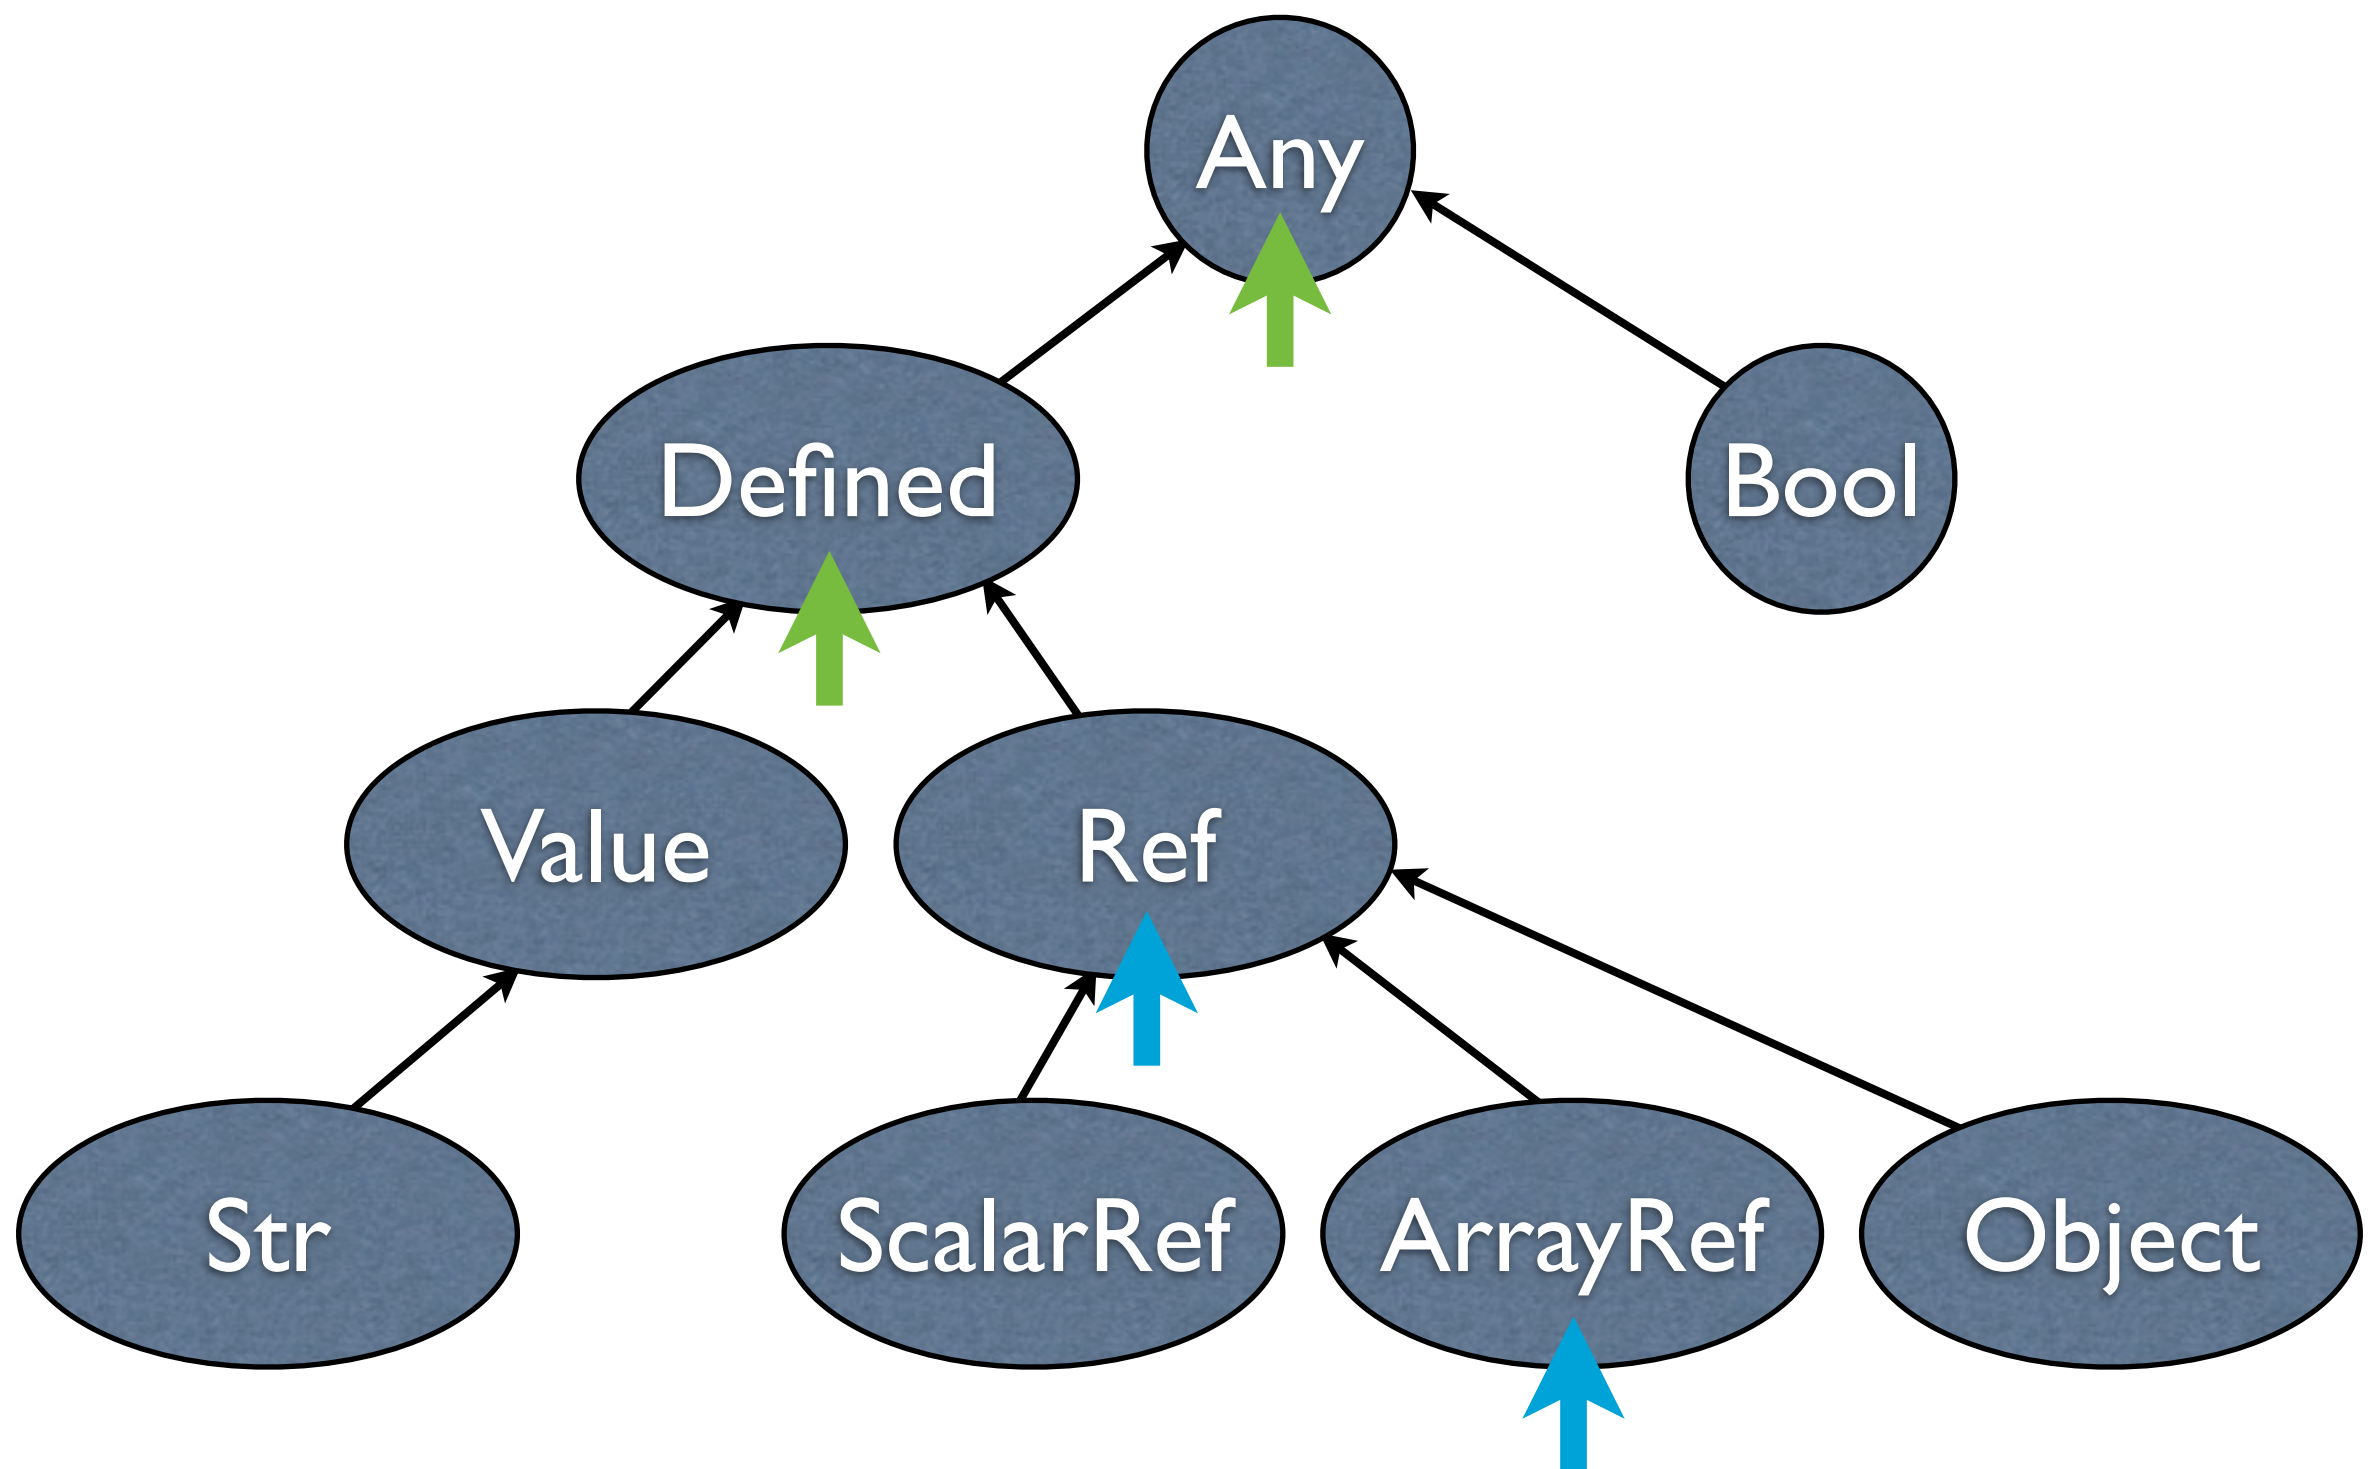


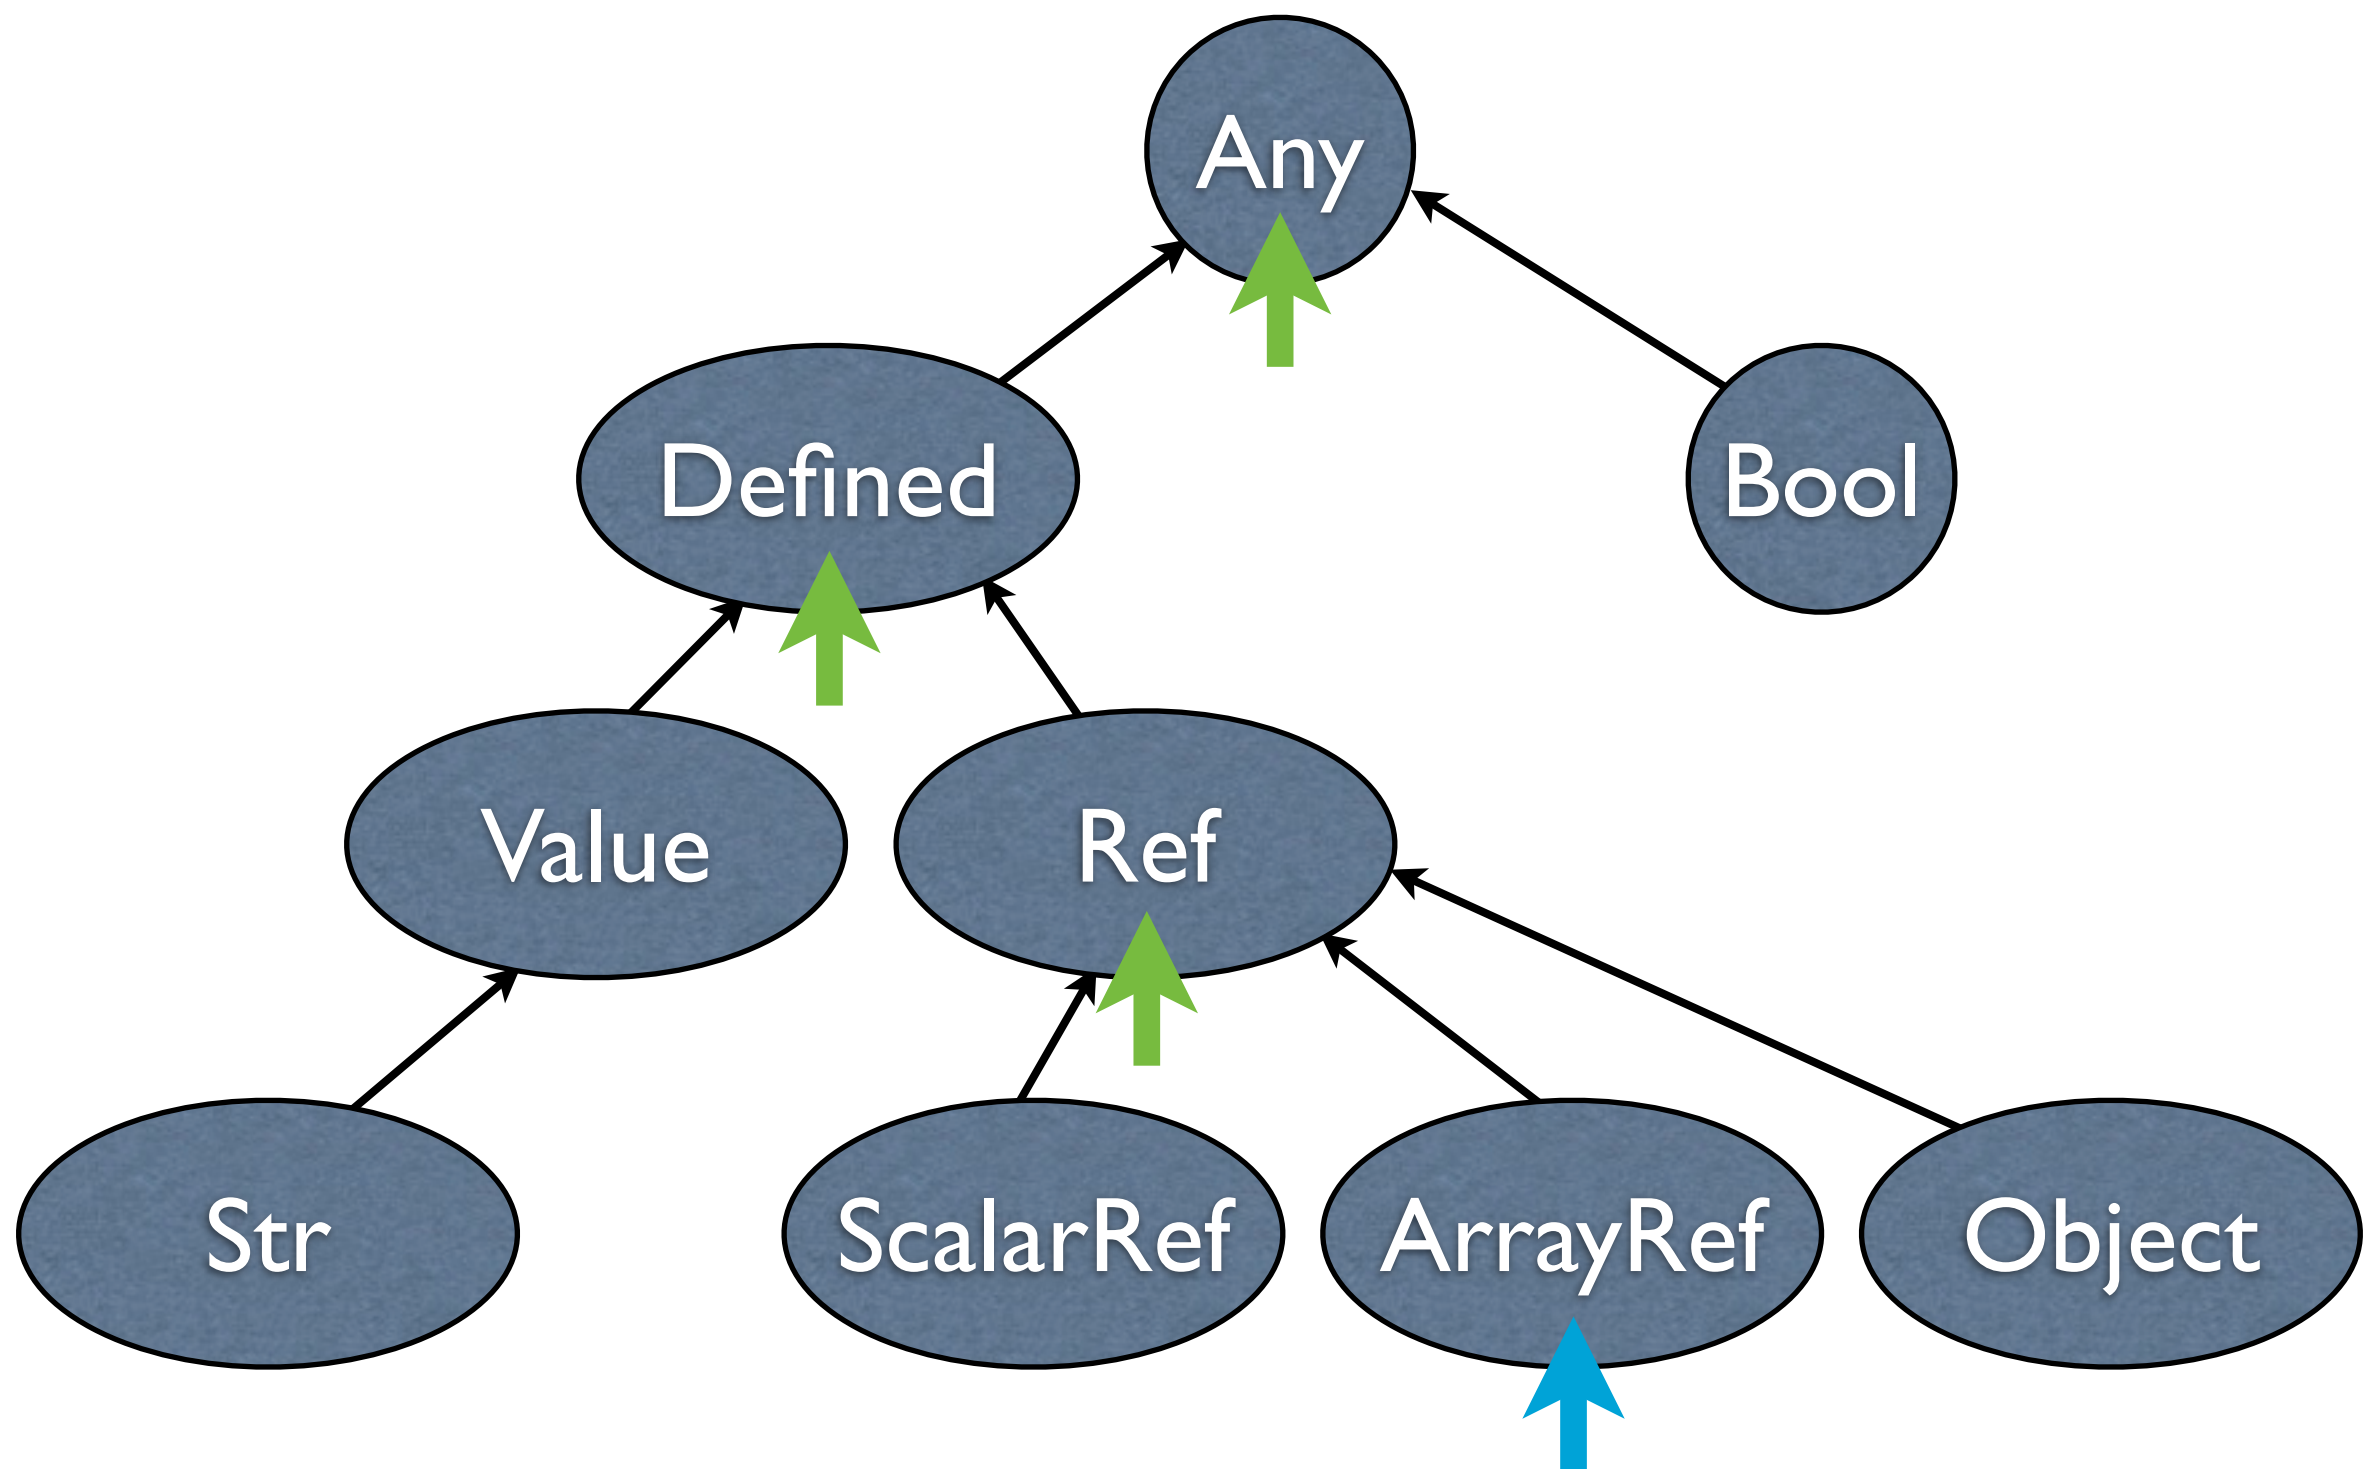


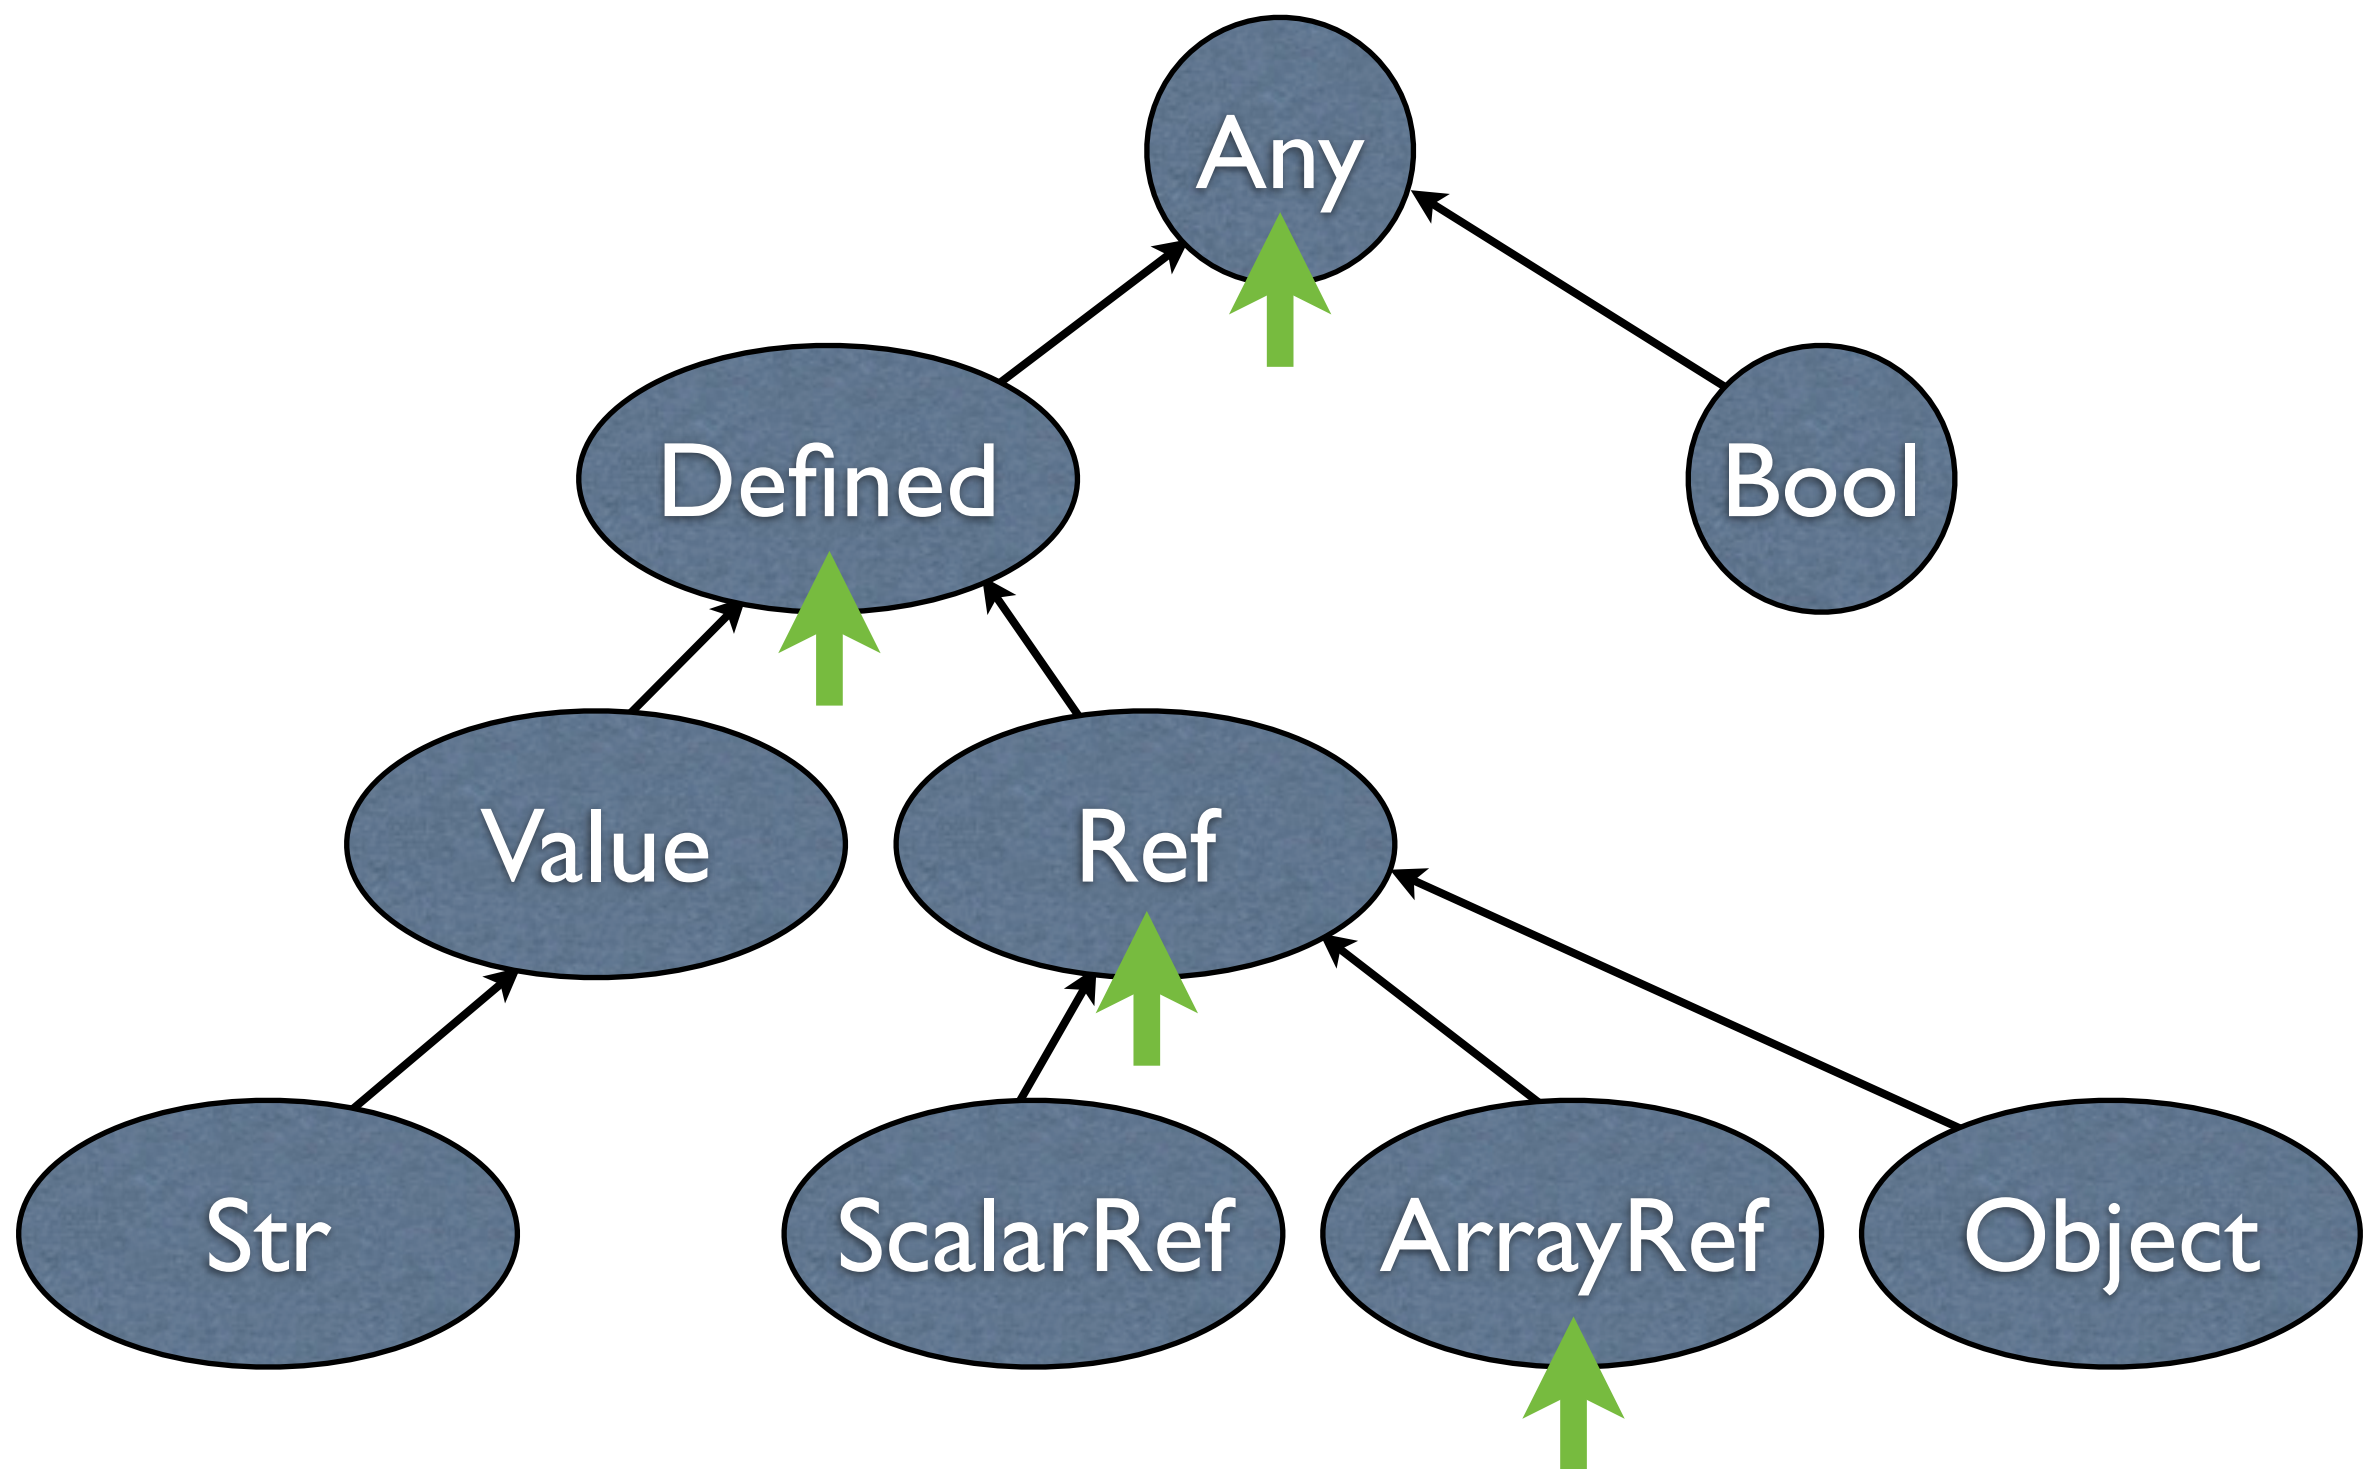












# Types add up!

- { 1 } # Any
- { defined(\$\_) } # Defined
- { ref(\$\_) } # Ref
- { ref(\$\_) eq 'ARRAY' } # ArrayRef

# Types add up!

- { 1 } # Any
- { defined(\$\_) } # Defined
- { ref(\$\_) } # Ref
- { ref(\$\_) eq 'ARRAY' } # ArrayRef

# inline\_as

- ▶ Optimize your type constraint check!
- ▶ Your parent constraints will not be called
- ▶ You give Moose a string, not a subroutine
- ▶ Formerly “optimized\_as”



# Types add up!

- { 1 } # Any
- { defined(\$\_) } # Defined
- { ref(\$\_) } # Ref
- { ref(\$\_) eq 'ARRAY' } # ArrayRef



# Type validation

```
subtype 'PositiveInt'
```

```
as 'Int',
```

```
where { $_ > 0 },
```

```
inline_as {
```

```
    $_[1] . ' && ' .
```

```
    $_[1] . '=~ /\d+$/'
```

```
};
```

# inline\_as: Our free lunch?

# Recap

- ▶ First rule of optimization: don't!
- ▶ Second rule of optimization: profile
- ▶ Devel::NYTProf, DTrace
- ▶ Fix your data structures and algorithms
- ▶ Can you cache?

# Recap

- ▶ `make_immutable` (almost) always
- ▶ Don't overuse attributes
- ▶ Role overhead is compile-time
- ▶ Lazy
- ▶ Types are expensive...
- ▶ ...so use `inline_as` to make them fast